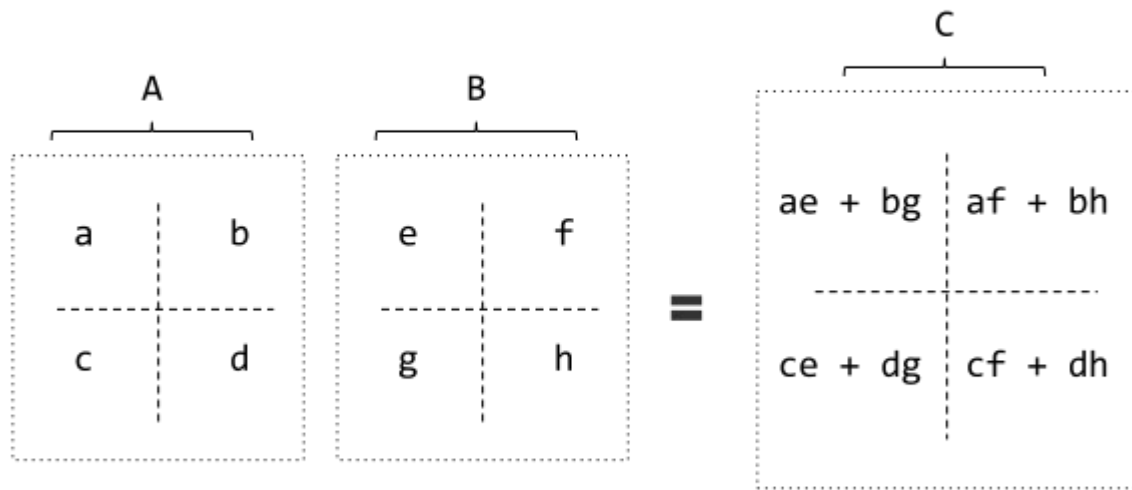

Optimization of Strassen's Method for Matrix Multiply

— Austin Schiller and Larry Sun —

Intuitive Recursive Matrix Multiply

DIVIDE AND CONQUER

- Divide both the A and B matrices into a **4** sub matrices
- Treating the submatrices as elements, do matrix multiplication to get the resulting C matrix
- Repeat for all **8** new matrix multiplies (ae, bg, af... ect.)
- Algorithm runs at $O(N^3)$



- A, B, and C are of size $N \times N$.
- $a, b, c, d, e, f, g,$ and h submatrices of size $N/2 \times N/2$

Strassen's Matrix Multiply

- Divide both the A and B matrices into a **4** sub matrices
- Create the **7** listed products using the submatrices (M1-M7)
- For all **7** submatrices M1-M7, do Strassen's Method on each submatrix!
- Once the submatrix is at a certain optimal size (size<=THRESHOLD), do regular matrix-matrix multiply
- Algorithm runs at $O(N^{2.8})$

$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}, B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix}, C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix}$$

$$C = \begin{bmatrix} M_1 + M_4 - M_5 + M_7 & M_3 + M_5 \\ M_2 + M_4 & M_1 + M_3 - M_2 + M_6 \end{bmatrix}$$

$$M_1 = (A_{11} + A_{22}) \times (B_{11} + B_{22})$$

$$M_2 = (A_{21} + A_{22}) \times B_{11}$$

$$M_3 = A_{11} \times (B_{12} + B_{22})$$

$$M_4 = A_{22} \times (B_{21} + B_{11})$$

$$M_5 = (A_{11} + A_{12}) \times B_{22}$$

$$M_6 = (A_{21} - A_{11}) \times (B_{11} + B_{12})$$

$$M_7 = (A_{12} - A_{22}) \times (B_{21} + B_{22})$$

- A, B, and C are of size NxN and the size is a power of 2
- a11, a12, a21, a22 are submatrices of size N/2 x N/2. Same with the b11, b12...c22

Why Optimizing Strassen's Method is Difficult

- Memory management-
 - How to do all of the calculations with minimal extra space (allocating submatrices, temp matrices, etc.)
 - When to stop recursively creating matrices and use the serial method
 - The correct order of operation to avoid data dependencies
- Strassen-Winograd method
 - Use $\frac{2}{3} N^2$ extra space with specific scheduling to further optimize the result

Strassen-Winograd method scheduling

#	operation	loc.	#	operation	loc.
1	$Z_1 = C_{22} - C_{12}$	C_{22}	14	$P_2 = \text{Acc}(\alpha A_{12} B_{21} + \beta C_{11})$	C_{11}
2	$Z_3 = C_{12} - C_{21}$	C_{12}	15	$U_1 = P_1 + P_2$	C_{11}
3	$S_1 = A_{21} + A_{22}$	X	16	$U_5 = U_2 + P_3$	C_{12}
4	$T_1 = B_{12} - B_{11}$	Y	17	$S_3 = A_{11} - A_{21}$	X
5	$P_5 = \text{Acc}(\alpha S_1 T_1 + \beta Z_3)$	C_{12}	18	$T_3 = B_{22} - B_{12}$	Y
6	$S_2 = S_1 - A_{11}$	X	19	$U_3 = P_7 + U_2$	C_{21}
7	$T_2 = B_{22} - T_1$	Y		$= \alpha \text{AccLR}(S_3 T_3 + U_2)$	
8	$P_6 = \text{Acc}(\alpha S_2 T_2 + \beta C_{21})$	C_{21}	20	$U_7 = U_3 + W_1$	C_{22}
9	$S_4 = A_{12} - S_2$	X	21	$T'_1 = B_{12} - B_{11}$	Y
10	$W_1 = P_5 + \beta Z_1$	C_{22}	22	$T'_2 = B_{22} - T'_1$	Y
11	$P_3 = \text{Acc}(\alpha S_4 B_{22} + P_5)$	C_{12}	23	$T_4 = T'_2 - B_{21}$	Y
12	$P_1 = \alpha A_{11} B_{11}$	X	24	$U_6 = U_3 - P_4$	C_{21}
13	$U_2 = P_6 + P_1$	C_{21}		$= -\alpha \text{AccR}(A_{22} T_4 - U_3)$	

Strassen Optimization: Saving Some Space

- Create new matrix add/sub functions that take in two **full** matrices and puts the result in a **smaller** submatrix

- **startrow, startcol** are passed in to specify which submatrix(a11, b21, etc.) of the **full** matrix are being worked on

- For example, if $\text{sublength} = \text{length}(A)/2$, then:

`m_add(A, B, c11, 0, 0, sublength, sublength)`

is the same as:

`a11+b22=c11`

- This saves unnecessary allocating of some submatrices (a11, a12, a21...)

$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}, B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix}, C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix}$$

```
void m_add(matrix_ptr a, matrix_ptr b, matrix_ptr c, int startrow1, int
startcol1, int startrow2, int startcol2)
{
    int i, j;
    int length = get_matrix_length(a);
    int sublength = get_matrix_length(c);
    data_t *m1 = get_matrix_start(a);
    data_t *m2 = get_matrix_start(b);
    data_t *m3 = get_matrix_start(c);

    for(i=0; i < sublength; i++){
        for(j=0; j < sublength; j++){
            m3[i*sublength + j] = m1[(i+startrow1)*length +
startcol1 + j] + m2[(i+startrow2)*length + startcol2 + j];
        }
    }
}
```

Strassen's Method with Loop Unrolling

- Unroll the arithmetic operations (matrix addition/subtraction)
- Unroll the data transferring of the main A, B matrices into the submatrices for work
- Unrolling greater than 8 was found to be ineffective...

$$\mathbf{A} = \begin{bmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{bmatrix}, \mathbf{B} = \begin{bmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{bmatrix}, \mathbf{C} = \begin{bmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{bmatrix}$$

```
for(i=0; i < sublength; i++){
    for(j=0; j < sublength; j+=UNROLL_THRESHOLD){
        a11->data[i*sublength+j] = a0[i*length+j];
        b11->data[i*sublength+j] = b0[i*length+j];
        a22->data[i*sublength+j] = a0[(i+sublength)*length+sublength+j];
        b22->data[i*sublength+j] = b0[(i+sublength)*length+sublength+j];

        a11->data[i*sublength+j+1] = a0[i*length+j+1];
        b11->data[i*sublength+j+1] = b0[i*length+j+1];
        a22->data[i*sublength+j+1] = a0[(i+sublength)
*length+sublength+j+1];
        b22->data[i*sublength+j+1] = b0[(i+sublength)
*length+sublength+j+1];

        a11->data[i*sublength+j+2] = a0[i*length+j+2];
        b11->data[i*sublength+j+2] = b0[i*length+j+2];
        a22->data[i*sublength+j+2] = a0[(i+sublength)
*length+sublength+j+2];
        b22->data[i*sublength+j+2] = b0[(i+sublength)
*length+sublength+j+2];

        a11->data[i*sublength+j+3] = a0[i*length+j+3];
        b11->data[i*sublength+j+3] = b0[i*length+j+3];
        a22->data[i*sublength+j+3] = a0[(i+sublength)
*length+sublength+j+3];
        b22->data[i*sublength+j+3] = b0[(i+sublength)
*length+sublength+j+3];
```

Strassen's Method with OpenMP

- Since we need to create 8 new submatrices (a11, a12...c22), we can parallelize the copying of the A, B, and C matrices to the submatrices
- Every time an operation is done on across elements of a submatrix, begin the loops with **parallel for** statements
- Matrix addition and matrix subtraction is omp optimized

$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}, B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix}, C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix} \quad C = \begin{bmatrix} M_1 + M_4 - M_5 + M_7 & M_3 + M_5 \\ M_2 + M_4 & M_1 + M_3 - M_2 + M_6 \end{bmatrix}$$

```
#pragma omp parallel shared(a0,b0,a11,a22,b11,b22,sublength,chunk) private(i,j) num_threads(OMP_THREADS)
{
    #pragma omp for schedule (dynamic,chunk) nowait
    for(i=0; i < sublength; ++i){
        for(j=0; j < sublength; ++j){
            a11->data[i*sublength+j] = a0[i*length+j];
            b11->data[i*sublength+j] = b0[i*length+j];
            a22->data[i*sublength+j] = a0[(i+sublength)
            *length+sublength+j];
            b22->data[i*sublength+j] = b0[(i+sublength)
            *length+sublength+j];
        }
        :
    }
    :
    main Strassen's work...
    :
}
```

```
#pragma omp parallel shared(c0,c11,c12,c21,c22,sublength,chunk) private(i,j) num_threads(OMP_THREADS)
{
    #pragma omp for schedule (dynamic,chunk) nowait
    for(i=0; i<sublength; ++i){
        for(j=0; j<sublength; ++j){
            c0[i*length+j] = c11->data[i*sublength+j];
            c0[i*length+sublength+j] = c12->data[i*sublength+j];
            c0[(i+sublength)*length+j] = c21->data[i*sublength+j];
            c0[(i+sublength)*length+sublength+j] = c22->data
            [i*sublength+j];
        }
    }
}
```

$$\begin{aligned} M_1 &= (A_{11} + A_{22}) \times (B_{11} + B_{22}) \\ M_2 &= (A_{21} + A_{22}) \times B_{11} \\ M_3 &= A_{11} \times (B_{12} + B_{22}) \\ M_4 &= A_{22} \times (B_{21} + B_{11}) \\ M_5 &= (A_{11} + A_{12}) \times B_{22} \\ M_6 &= (A_{21} - A_{11}) \times (B_{11} + B_{12}) \\ M_7 &= (A_{12} - A_{22}) \times (B_{21} + B_{22}) \end{aligned}$$

Strassen's Using Pthreads

- Create 7 threads, each thread works on one of the M_i matrices, where i ranges from 1 to 7
- Call the recursive function inside each of the 7 threads
- When a small matrix size is reached, finish with the original matrix multiply function
- When all threads are finished with their M_i , join them together in main function and do the arithmetic to get C

$$A = \begin{bmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{bmatrix}, B = \begin{bmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{bmatrix}, C = \begin{bmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{bmatrix}$$

The diagram illustrates the parallel execution of Strassen's algorithm using 7 threads (P1 through P7). Each thread calculates a specific intermediate matrix M_i based on the input matrices A and B . The threads are arranged vertically, and arrows point from each thread to its corresponding M_i equation. A line from the fourth bullet point in the list above points to the threads, indicating that the main function joins the threads and performs the final arithmetic to produce C .

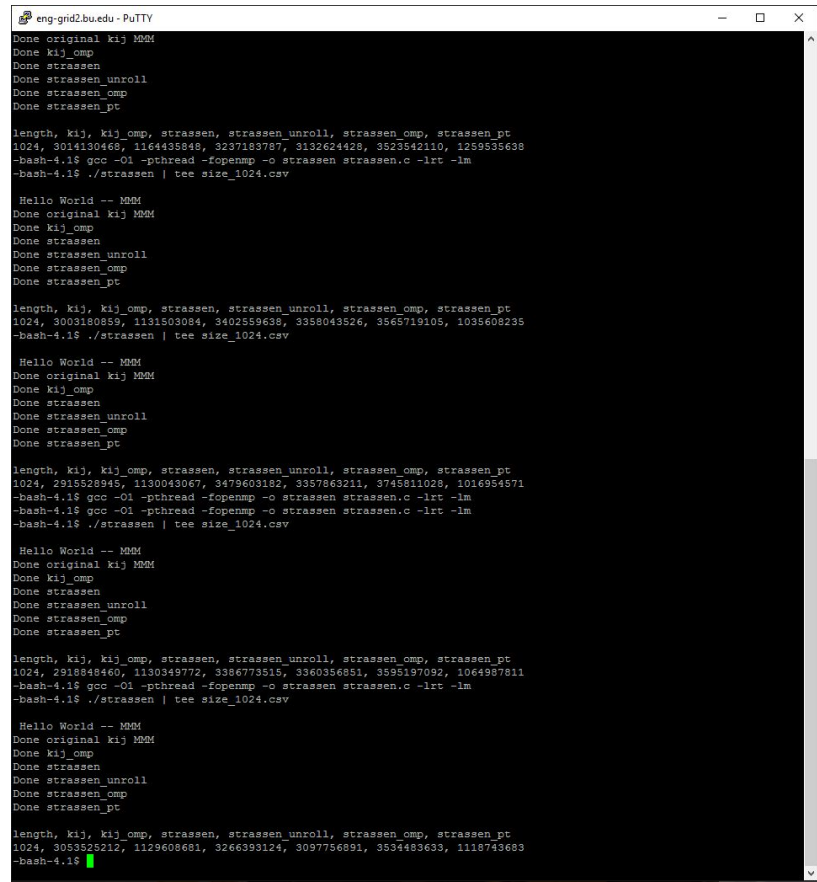
$$C = \begin{bmatrix} M_1 + M_4 - M_5 + M_7 & M_3 + M_5 \\ M_2 + M_4 & M_1 + M_3 - M_2 + M_6 \end{bmatrix}$$

Thread assignments:

- P1: $M_1 = (A_{11} + A_{22}) \times (B_{11} + B_{22})$
- P2: $M_2 = (A_{21} + A_{22}) \times B_{11}$
- P3: $M_3 = A_{11} \times (B_{12} + B_{22})$
- P4: $M_4 = A_{22} \times (B_{21} + B_{11})$
- P5: $M_5 = (A_{11} + A_{12}) \times B_{22}$
- P6: $M_6 = (A_{21} - A_{11}) \times (B_{11} + B_{12})$
- P7: $M_7 = (A_{12} - A_{22}) \times (B_{21} + B_{22})$

Experimental Procedure

- Tested on eng-grid2
- Used -O1 optimization
- Since Strassen's Method performs worse than serial MMM at **small** matrix sizes (due to recursion/mem allocs), we tested with sizes: **512, 1024, 2048, 4096**
- In the Strassen versions, the optimal threshold for starting normal MMM is different at different matrix sizes.
(On average: Threshold = length/8 or 16)



```
eng-grid2.bu.edu - PuTTY
Done original ksj MMM
Done ksj_omp
Done strassen
Done strassen_unroll
Done strassen_omp
Done strassen_pt

length, ksj, ksj_omp, strassen, strassen_unroll, strassen_omp, strassen_pt
1024, 3014130468, 1164495946, 3237183787, 3132624428, 3523542110, 1269535638
-bash-4.1$ gcc -O1 -pthread -fopenmp -o strassen strassen.c -lrt -lm
-bash-4.1$ ./strassen | tee size_1024.csv

Hello World -- MMM
Done original ksj MMM
Done ksj_omp
Done strassen
Done strassen_unroll
Done strassen_omp
Done strassen_pt

length, ksj, ksj_omp, strassen, strassen_unroll, strassen_omp, strassen_pt
1024, 3003180859, 1131503084, 3402559638, 3358043526, 3565719105, 1035608235
-bash-4.1$ ./strassen | tee size_1024.csv

Hello World -- MMM
Done original ksj MMM
Done ksj_omp
Done strassen
Done strassen_unroll
Done strassen_omp
Done strassen_pt

length, ksj, ksj_omp, strassen, strassen_unroll, strassen_omp, strassen_pt
1024, 2915528945, 1130043067, 3479603182, 3357863211, 3745811028, 1016954571
-bash-4.1$ gcc -O1 -pthread -fopenmp -o strassen strassen.c -lrt -lm
-bash-4.1$ ./strassen | tee size_1024.csv

Hello World -- MMM
Done original ksj MMM
Done ksj_omp
Done strassen
Done strassen_unroll
Done strassen_omp
Done strassen_pt

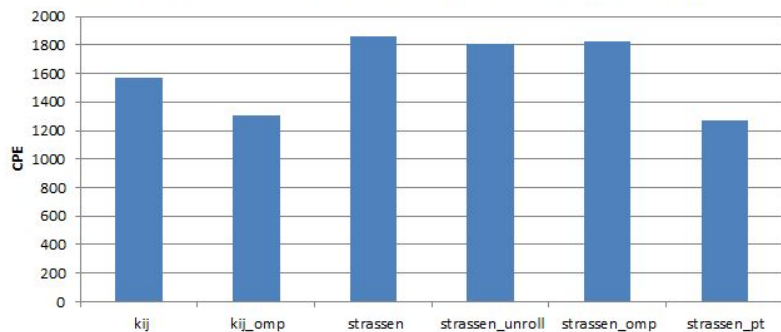
length, ksj, ksj_omp, strassen, strassen_unroll, strassen_omp, strassen_pt
1024, 2918848460, 1130349772, 3386773515, 3360356851, 3595197092, 1064987811
-bash-4.1$ gcc -O1 -pthread -fopenmp -o strassen strassen.c -lrt -lm
-bash-4.1$ ./strassen | tee size_1024.csv

Hello World -- MMM
Done original ksj MMM
Done ksj_omp
Done strassen
Done strassen_unroll
Done strassen_omp
Done strassen_pt


length, ksj, ksj_omp, strassen, strassen_unroll, strassen_omp, strassen_pt
1024, 3053525212, 1129608681, 3266393124, 3097756891, 3534483633, 1118743683
-bash-4.1$
```

Results


MMM vs Strassen's Methods (Matrix Length: 512)



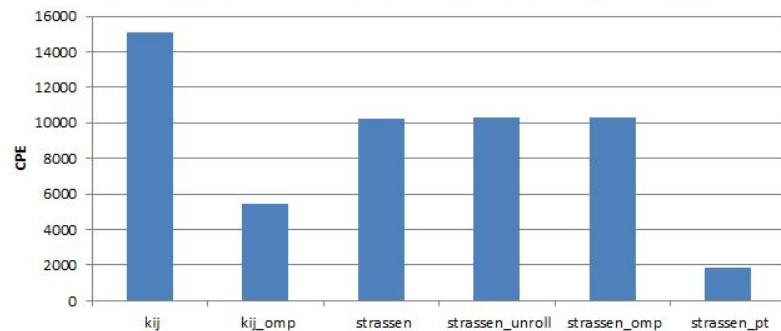
MMM vs Strassen's Methods (Matrix Length: 1024)



MMM vs Strassen's Methods (Matrix Length: 2048)



MMM vs Strassen's Methods (Matrix Length: 4096)



Conclusion

- At large matrix sizes, the Strassen methods perform significantly better
- Strassen omp and unrolled versions performs about the same as normal version
- Strassen with Pthreads performs very well! (At 4096, almost a factor of 8x speedup)
- MMM_omp still performs better than most Strassen versions, but at the cost of some error (~1%) due to race conditions

Possible Improvements in the Future

- Perform more fine-grained omp sections in the Strassen version (perhaps assign the 7 matrices to individual threads like in pthreads)
- Unroll the recursion and implement iteration to allow for more parallelism
- Implement Strassen's Method with Intel Intrinsics (perform vector arithmetic with the smaller submatrices?)
- Maybe implement blocking
- Combine the best methods

Thank You!